

data-loading-feature-extraction

January 10, 2026

1 Mount Google Drive

```
[ ]: # Google Colab specific integration
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

2 Install Required Packages

```
[ ]: print(" Installing required packages...")

!pip install -q torch torchvision
!pip install -q scikit-learn
!pip install -q umap-learn
!pip install -q pillow
!pip install -q matplotlib seaborn
!pip install -q tqdm
!pip install -q pyyaml
!pip install rarfile

print(" All packages installed successfully!")
```

```
Installing required packages...
23.8/23.8 MB
118.2 MB/s eta 0:00:00
Collecting rarfile
  Downloading rarfile-4.2-py3-none-any.whl.metadata (4.4 kB)
  Downloading rarfile-4.2-py3-none-any.whl (29 kB)
Installing collected packages: rarfile
Successfully installed rarfile-4.2
All packages installed successfully!
```

3 Create project root in Google Drive

```
[ ]: import os
PROJECT_ROOT = '/content/drive/MyDrive/image-clustering-retrieval'
os.makedirs(PROJECT_ROOT, exist_ok=True)

# Create folder structure
folders = [
    'data/raw',
    'data/processed',
    'features',
    'models/pretrained',
    'models/clustering',
    'models/classifier',
    'models/faiss_index',
    'results/clustering',
    'results/retrieval/sample_queries',
    'results/classification',
    'configs'
]

for folder in folders:
    os.makedirs(f'{PROJECT_ROOT}/{folder}', exist_ok=True)

print(" Project structure created!")
print(f" Project root: {PROJECT_ROOT}")

# Change working directory
os.chdir(PROJECT_ROOT)
print(f" Current directory: {os.getcwd()}")
```

```
Project structure created!
Project root: /content/drive/MyDrive/image-clustering-retrieval
Current directory: /content/drive/MyDrive/image-clustering-retrieval
```

```
[ ]: # Set random seeds
import torch
import numpy as np
torch.manual_seed(42)
np.random.seed(42)
```

```
[ ]: # check for GPU

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f" Using device: {device}")
if device.type == 'cuda':
    print(f"   GPU: {torch.cuda.get_device_name(0)}")
```

```
    print(f"    Memory: {torch.cuda.get_device_properties(0).total_memory / 1e9:.  
        <2f} GB")
```

```
Using device: cuda  
GPU: Tesla T4  
Memory: 15.83 GB
```

4 Extract Caltech-101 Dataset from RAR (Colab Local Storage)

```
[ ]: import os  
import rarfile  
  
print("==" * 60)  
print("EXTRACTING CALTECH-101 DATASET INTO PROCESSED FOLDER (GOOGLE DRIVE)")  
print("==" * 60)  
  
# Install RAR support (run once per runtime)  
!apt-get -qq install unrar  
!pip install -q rarfile  
  
# Path to uploaded RAR file (RAW source)  
uploaded_rar_path = "/content/drive/MyDrive/image-clustering-retrieval/data/raw/  
object.rar"  
  
# Target directory (PROCESSED)  
processed_dir = "/content/drive/MyDrive/image-clustering-retrieval/data/  
processed"  
dataset_path = os.path.join(processed_dir, "object")  
  
# Sanity check: RAR exists  
if not os.path.exists(uploaded_rar_path):  
    raise FileNotFoundError(  
        f" RAR file not found at {uploaded_rar_path}"  
    )  
  
# Create processed directory if missing  
os.makedirs(processed_dir, exist_ok=True)  
  
# Extract only if not already extracted  
if not os.path.exists(dataset_path):  
    print(" Extracting dataset into processed folder...")  
  
    with rarfile.RarFile(uploaded_rar_path) as rf:  
        rf.extractall(processed_dir)  
  
    print(" Extraction complete!")
```

```

else:
    print(" Dataset already exists - skipping extraction")

# Verify dataset structure
if not os.path.exists(dataset_path):
    raise RuntimeError(" Extraction failed: 'object/' folder not found!")

# Dataset stats
categories = sorted([
    d for d in os.listdir(dataset_path)
    if os.path.isdir(os.path.join(dataset_path, d)) and not d.startswith("."))
])

total_images = sum(
    len([
        f for f in os.listdir(os.path.join(dataset_path, c))
        if f.lower().endswith((".jpg", ".jpeg", ".png")))
    ])
    for c in categories
)

print("\n Dataset Statistics:")
print(f" Total categories: {len(categories)}")
print(f" Total images: {total_images}")
print(f"\n Dataset location (PROCESSED): {dataset_path}")

```

```
=====
EXTRACTING CALTECH-101 DATASET INTO PROCESSED FOLDER (GOOGLE DRIVE)
=====
```

```

Extracting dataset into processed folder...
Extraction complete!

```

```

Dataset Statistics:
    Total categories: 50
    Total images: 4696

```

```

Dataset location (PROCESSED): /content/drive/MyDrive/image-clustering-
retrieval/data/processed/object

```

5 Create Custom Dataset Class

```
[ ]: import os
from torch.utils.data import Dataset
from PIL import Image
from tqdm import tqdm

class ObjectDataset(Dataset):

```

```

"""
Custom Dataset for Caltech-style object dataset
Folder structure:
/content/object/
    accordion/
    airplanes/
    ...
"""

def __init__(self, root_dir="/content/object", transform=None, ↴
             max_images_per_class=None):
    self.root_dir = root_dir
    self.transform = transform
    self.max_images_per_class = max_images_per_class

    # Get class folders
    self.categories = sorted([
        d for d in os.listdir(root_dir)
        if os.path.isdir(os.path.join(root_dir, d))
    ])

    self.category_to_idx = {cat: idx for idx, cat in enumerate(self. ↴
                                                               categories)}

    self.image_paths = []
    self.labels = []

    print(f" Found {len(self.categories)} categories")

    for category in tqdm(self.categories, desc="Scanning categories"):
        cat_path = os.path.join(root_dir, category)
        cat_idx = self.category_to_idx[category]

        images = [
            f for f in os.listdir(cat_path)
            if f.lower().endswith('.jpg', '.jpeg', '.png'))
        ]

        if max_images_per_class:
            images = images[:max_images_per_class]

        for img in images:
            self.image_paths.append(os.path.join(cat_path, img))
            self.labels.append(cat_idx)

    print(f" Loaded {len(self.image_paths)} images")

```

```

def __len__(self):
    return len(self.image_paths)

def __getitem__(self, idx):
    img_path = self.image_paths[idx]
    label = self.labels[idx]

    try:
        image = Image.open(img_path).convert("RGB")
    except:
        image = Image.new("RGB", (224, 224), color="black")

    if self.transform:
        image = self.transform(image)

    return image, label

def get_category_name(self, idx):
    return self.categories[idx]

```

6 PHASE 1: DATA PREPARATION”

```

[ ]: import torchvision.transforms as transforms
import json
from datetime import datetime
import matplotlib.pyplot as plt

print("-"*60)
print("PHASE 1: DATA PREPARATION")
print("-"*60)

# Define transforms
transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(
        mean=[0.485, 0.456, 0.406],
        std=[0.229, 0.224, 0.225]
    )
])

# Use CORRECT dataset class
full_dataset = ObjectDataset(
    root_dir=dataset_path,    # /content/object
    transform=transform,
    max_images_per_class=None
)

```

```

)

print(f"\n Dataset created:")
print(f" Total images: {len(full_dataset)}")
print(f" Total categories: {len(full_dataset.categories)}")

# Save metadata
metadata = {
    "num_categories": len(full_dataset.categories),
    "num_images": len(full_dataset),
    "categories": full_dataset.categories,
    "category_to_idx": full_dataset.category_to_idx,
    "dataset_path": dataset_path,
    "created_at": datetime.now().strftime("%Y-%m-%d %H:%M:%S")
}

os.makedirs(f"{PROJECT_ROOT}/configs", exist_ok=True)

with open(f"{PROJECT_ROOT}/configs/dataset_metadata.json", "w") as f:
    json.dump(metadata, f, indent=2)
print(" Metadata saved")

# Visualize samples
def visualize_samples(dataset, num_samples=10):
    fig, axes = plt.subplots(2, 5, figsize=(15, 6))
    axes = axes.ravel()

    indices = np.random.choice(len(dataset), num_samples, replace=False)

    mean = np.array([0.485, 0.456, 0.406])
    std = np.array([0.229, 0.224, 0.225])

    for i, ax in enumerate(axes):
        img, label = dataset[indices[i]]

        img = img.numpy().transpose(1, 2, 0)
        img = std * img + mean
        img = np.clip(img, 0, 1)

        ax.imshow(img)
        ax.set_title(dataset.get_category_name(label), fontsize=9)
        ax.axis("off")

    plt.tight_layout()
    plt.show()

print("\n Visualizing sample images...")

```

```
visualize_samples(full_dataset)
```

PHASE 1: DATA PREPARATION

Found 50 categories

Scanning categories: 100% | 50/50 [00:00<00:00, 577.00it/s]

Loaded 4696 images

Dataset created:

Total images: 4696

Total categories: 50

Metadata saved

Visualizing sample images...



7 PHASE 2: FEATURE EXTRACTION

7.1 Feature Extraction with ResNet-50

```
[ ]: import torch
import torch.nn as nn
import torchvision.models as models
from torch.utils.data import DataLoader, Dataset
import numpy as np
from sklearn.preprocessing import normalize
from tqdm import tqdm
import os
from PIL import Image
```

```

import torchvision.transforms as transforms

print("-"*60)
print("PHASE 2: FEATURE EXTRACTION")
print("-"*60)

# --- Start: Re-definitions for kernel stability ---
# This block ensures essential variables are defined even after a kernel
# restart.
# Ideally, previous cells (byqToj4MW1nX, zlz_pNMtYJYY, uL1P5rKukEm9,
# F6ZHVLZ0kG1j, YB09S1EiXayz)
# should be executed first.

# Define PROJECT_ROOT if not already in scope (from byqToj4MW1nX)
if 'PROJECT_ROOT' not in globals():
    PROJECT_ROOT = '/content/drive/MyDrive/image-clustering-retrieval'
    os.makedirs(PROJECT_ROOT, exist_ok=True)
    print(f'Re-defined PROJECT_ROOT: {PROJECT_ROOT}')

# Define dataset_path if not already in scope (from zlz_pNMtYJYY)
if 'dataset_path' not in globals():
    processed_dir = os.path.join(PROJECT_ROOT, "data/processed")
    dataset_path = os.path.join(processed_dir, "object")
    os.makedirs(processed_dir, exist_ok=True)
    print(f'Re-defined dataset_path: {dataset_path}')

# Define ObjectDataset class if not already in scope (from uL1P5rKukEm9)
if 'ObjectDataset' not in globals():
    print("Re-defining ObjectDataset class...")
    class ObjectDataset(Dataset):
        """
        Custom Dataset for Caltech-style object dataset
        Folder structure:
        /content/object/
            accordion/
            airplanes/
            ...
        """
        def __init__(self, root_dir, transform=None, max_images_per_class=None):
            self.root_dir = root_dir
            self.transform = transform
            self.max_images_per_class = max_images_per_class
            self.categories = sorted([d for d in os.listdir(root_dir) if os.
path.isdir(os.path.join(root_dir, d))])
            self.category_to_idx = {cat: idx for idx, cat in enumerate(self.
categories)}
            self.image_paths = []

```

```

        self.labels = []
        for category in tqdm(self.categories, desc="Scanning categories\u2193(re-init)", leave=False):
            cat_path = os.path.join(root_dir, category)
            cat_idx = self.category_to_idx[category]
            images = [f for f in os.listdir(cat_path) if f.lower().endswith('.jpg', '.jpeg', '.png'))]
            if max_images_per_class:
                images = images[:max_images_per_class]
            for img in images:
                self.image_paths.append(os.path.join(cat_path, img))
                self.labels.append(cat_idx)
    def __len__(self):
        return len(self.image_paths)
    def __getitem__(self, idx):
        img_path = self.image_paths[idx]
        label = self.labels[idx]
        try:
            image = Image.open(img_path).convert("RGB")
        except:
            image = Image.new("RGB", (224, 224), color="black")
        if self.transform:
            image = self.transform(image)
        return image, label
    def get_category_name(self, idx):
        return self.categories[idx]

# Define transform if not already in scope (from F6ZHVLZ0kG1j)
if 'transform' not in globals():
    print("Re-defining transform...")
    transform = transforms.Compose([
        transforms.Resize((224, 224)),
        transforms.ToTensor(),
        transforms.Normalize(
            mean=[0.485, 0.456, 0.406],
            std=[0.229, 0.224, 0.225]
        )
    ])
else:
    print("Transform already defined")

# Define full_dataset if not already in scope (from F6ZHVLZ0kG1j)
if 'full_dataset' not in globals():
    print("Re-initializing full_dataset...")
    full_dataset = ObjectDataset(
        root_dir=dataset_path,
        transform=transform,
        max_images_per_class=None
    )

```

```

    print(f"    full_dataset re-initialized with {len(full_dataset)} images.")

# Define device if not already in scope (from YB09S1EiXayz)
if 'device' not in globals():
    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    print(f"Re-defined device: {device}")
# --- End: Re-definitions for kernel stability ---

class FeatureExtractor:
    def __init__(self, model_name='resnet50', device='cuda'):
        self.device = torch.device(device if torch.cuda.is_available() else
                                   'cpu')
        self.model_name = model_name

    print(f"\n Loading pretrained {model_name}...")

    if model_name == 'resnet50':
        model = models.resnet50(weights=models.ResNet50_Weights.
                               IMAGENET1K_V1)
        self.feature_dim = 2048
    else:
        raise ValueError("Only resnet50 supported")

    self.model = nn.Sequential(*list(model.children())[:-1])
    self.model.to(self.device)
    self.model.eval()

    print(f" Model loaded on {self.device}")
    print(f" Feature dimension: {self.feature_dim}")

def extract_features(self, dataloader, desc="Extracting features"):
    features_list = []
    labels_list = []

    with torch.no_grad():
        for images, labels in tqdm(dataloader, desc=desc):
            images = images.to(self.device)

            features = self.model(images)
            features = features.view(features.size(0), -1)

            features_list.append(features.cpu().numpy())
            labels_list.append(labels.numpy())

    features = np.vstack(features_list)
    labels = np.concatenate(labels_list)

```

```

        return features, labels
# Dataloader
batch_size = 64
dataloader = DataLoader(
    full_dataset,
    batch_size=batch_size,
    shuffle=False,
    num_workers=2,
    pin_memory=True
)

# Feature extractor
extractor = FeatureExtractor(model_name='resnet50', device=device)

print(f"\n Extracting features from {len(full_dataset)} images...")

features, labels = extractor.extract_features(dataloader)

# Normalize
features_normalized = normalize(features, norm='l2')

# Save
os.makedirs(f"{PROJECT_ROOT}/features", exist_ok=True)

np.save(f"{PROJECT_ROOT}/features/embeddings.npy", features_normalized)
np.save(f"{PROJECT_ROOT}/features/labels.npy", labels)
print(" Feature extraction & saving complete!")

```

PHASE 2: FEATURE EXTRACTION

Loading pretrained resnet50...
 Model loaded on cuda
 Feature dimension: 2048

Extracting features from 4696 images...

Extracting features: 100% | 74/74 [00:29<00:00, 2.55it/s]

Feature extraction & saving complete!

[]: