

# MACHINE LEARNING

## Complex Computing Activity

### PROJECT REPORT

**Image Clustering and Retrieval System**

#### **Team Members**

Abdul Basit (22F-BSAI-25)

Shehryar Ahmed (22F-BSAI-28)

Muhammad Hamamd (22F-BSAI-39)

Muhammad Sohaib (22F-BSAI-40)

**Objective:** To extract discriminative image features using a pretrained deep CNN.  
To compare clustering and classification algorithms for image organization and prediction.  
To build an interactive image retrieval system with measurable performance metrics.

#### **Project Description:**


- This project implements an image clustering, retrieval, and classification system using deep feature embeddings.
- Images from the Caltech-101 dataset are processed through a pretrained ResNet-50 model to extract high-level visual features.
- Unsupervised learning techniques (K-Means and Hierarchical clustering) group visually similar images.
- Supervised classifiers (SVM and Random Forest) are trained on extracted features for category prediction.
- A Streamlit application enables users to upload an image and retrieve visually similar images with confidence scores.

# 1. Introduction:

- With the rapid growth of visual data, automatic image understanding has become essential.
- Traditional image processing methods struggle to capture high-level semantic information.
- Deep learning-based feature extraction combined with classical machine learning provides a practical solution.
- This project demonstrates an end-to-end pipeline for image clustering, classification, and retrieval using real datasets.

## 2. Architecture & Pipeline

### 2.1 System Architecture



### Components:

- Data Input:** Images from the Caltech-101 dataset are loaded from disk using a structured directory-based dataset loader.
- Preprocessing Module:** Images are resized, normalized, and passed through a pretrained CNN to extract deep feature embeddings.
- Feature Extraction:** A pretrained ResNet-50 model generates 2048-dimensional feature vectors for each image.
- Feature Normalization:** Extracted features are L2-normalized to ensure distance-based algorithms behave correctly.
- Clustering Model:** Hirarchical (K=50) groups images into visually similar clusters in an unsupervised manner.
- Clustering Evaluation:** Silhouette Score and Davies–Bouldin Index are used to assess cluster quality.

7. **Dimensionality Reduction:** Optional PCA reduces feature dimensions before classification to improve efficiency.
8. **Classification Model:** A Support Vector Machine classifier is trained on the extracted features for supervised category prediction.
9. **Evaluation Module:** Classification accuracy and confusion matrices measure model performance.
10. **Output/Results:** The system produces cluster assignments, trained clustering and classification models, evaluation metrics, and visualizations (t-SNE, confusion matrix, performance plots). A Streamlit application enables user interaction by allowing image upload, retrieving visually similar images from the dataset, and displaying similarity scores/probabilities with image names in real time.

## 2.2 Pipeline Diagram

```
[Caltech-101 Images]
↓
[Image Loading]
↓
[Resize + Normalize]
↓
[ResNet-50 Feature Extraction]
↓
[2048-D Feature Vectors]
↓
[L2 Normalization]
↓
[Hierarchical Agglomerative Clustering (K=50) ]
↓
[Silhouette & Davies-Bouldin Evaluation]
↓
[SVM Classification]
↓
[Accuracy + Confusion Matrix]
↓
[Results: Clusters, Models, Metrics, t-SNE Plots, image retrieval]
```

## 3. Lab Concepts Integration

Lab	Lab Concept	How It Was Used	Where in Code / Pipeline
5	Random Forest Classifier	Used as a supervised image classification model trained on ResNet-50 feature embeddings to predict image categories and compared against SVM.	models/classifier/random_forest.pkl, training in <b>Phase 4: Classification</b> , results in classification_results.json, visualization in results/classification/
6	Support Vector Machine (SVM) Classifier	Trained as an alternative supervised classifier on the same deep features to compare performance with Random Forest.	models/classifier/svm_rbf.pkl, results in classification_results_svm.json, accuracy plot rf_vs_svm_accuracy.png
10	K-Means Clustering	Applied for unsupervised grouping of image embeddings to analyze visual similarity and cluster structure before retrieval.	models/clustering/kmeans_k20.pkl, labels in cluster_labels.npy, metrics in clustering_metrics.json, elbow plot in results/clustering/elbow_curve.png
11	Hierarchical Agglomerative Clustering	Used to create fine-grained visual clusters (K=50) for better semantic grouping and cluster evaluation.	models/clustering/hir_model_k50.pkl, labels in hir_labels_k50.npy, metrics in hir_metrics_k50.json, dendrogram/TSNE in results/clustering/

## 4. Dataset

### 4.1 Dataset Description

1. **Name:** Caltech-101
2. **Source:** Caltech Vision Lab (official dataset) Size: [Number of samples, file size]
3. **Structure:** ~5,163 images across 51 object categories
4. **Split:** 80% training / 20% validation (for ML)

### 4.2 Preprocessing Steps

1. **Image Resizing:** All images resized to **224×224** to match CNN input requirements
2. **Normalization:** Pixel values normalized using ImageNet mean and standard deviation
3. **Feature Extraction:** Images passed through pretrained **ResNet-50** to obtain deep embeddings
4. **Feature Normalization:** L2 normalization applied to extracted feature vectors
5. **Dimensionality Reduction:** Optional **PCA** applied before classification for efficiency

## 5. Model Design & Training

### 5.1 Model Architecture

**Model Type:** ResNet50 (Pretrained Model)

**Architecture Details:**

- **Input Layer:** RGB image, 224×224×3
- **Initial Conv Layer:** 7×7 convolution, 64 filters, stride 2
- **Residual Blocks:**
  - Conv2\_x → Conv5\_x (bottleneck blocks with skip connections)
- **Global Average Pooling:** Reduces spatial dimensions
- **Output Feature Layer:** 2048-dimensional feature vector
- **Classification Layer: Removed / not used** (model used as feature extractor only)

### 5.2 Hyperparameters

Since ResNet-50 is frozen, most training hyperparameters do not apply.

Hyperparameter	Value	Rationale
Learning Rate	N/A	No CNN fine-tuning performed
Batch Size	32	Efficient GPU/CPU feature extraction
Epochs	1	Single forward pass for embeddings
Optimizer	N/A	Model weights not updated
Loss Function	N/A	No CNN training

### 5.3 Training Details

- **Hardware Used:** CPU (local machine / laptop)
- **Training Time:** Not applicable (only feature extraction)


- **Framework:** PyTorch
- **Key Techniques:**
  - Transfer learning (feature extraction)
  - L2 feature normalization
  - PCA before SVM/RF classification

## 6. Evaluation & Results


### 6.1 Evaluation Metrics

- **Primary Metric (Classification Accuracy):** 85.19% (SVM validation accuracy)
- **Clustering Quality Metrics:**
  - **Silhouette Score:** 0.1409 → Low, indicates clusters have some overlap.
  - **Davies-Bouldin Index:** 2.2168 → Higher than ideal, confirming clusters are not very compact.
- **Retrieval Performance:**
  - **Precision@10:** 0.9853 → Extremely high, top 10 retrieved images are almost always relevant.

**Confusion Matrix:**



**Graphs (Cluster visualization):**



### Key Findings:

- **Classification is reliable:** The SVM model generalizes well with ~86% validation accuracy.
- **Clustering needs improvement:** Low silhouette score and high Davies-Bouldin indicate clusters overlap, so hierarchical clustering might not fully separate categories.
- **Retrieval is highly effective:** Despite mediocre clustering, FAISS-based image retrieval achieves near-perfect precision@10, showing feature embeddings are strong for similarity search.

## 7. Challenges & Limitations

### 7.1 Challenges Faced

1. **High-dimensional feature handling:**
  - o **Description:** ResNet-50 outputs 2048-dimensional feature vectors, which increased memory usage and computation time for clustering and retrieval.
  - o **Solution:** Used efficient data structures and batch processing to handle features without crashing the system.
2. **Suboptimal clustering quality:**
  - o **Description:** Hierarchical clustering produced low silhouette scores and high Davies-Bouldin index, making some clusters overlap.
  - o **Solution:** Experimented with different linkage methods and distance metrics, but dataset diversity still caused imperfect separation.
3. **Large retrieval index:**
  - o **Description:** FAISS index stored all 5163 vectors, which increased memory footprint.
  - o **Solution:** Optimized FAISS parameters and limited search to top-k neighbors for faster queries.

### 7.2 Limitations

- **Dataset size and diversity:** Only 51 categories and 5163 images limit generalization to broader datasets.
- **Clustering effectiveness:** Hierarchical clustering struggles with fine-grained category separation; more advanced methods like spectral clustering or deep clustering could perform better.
- **Computational requirements:** Training SVM on high-dimensional embeddings and building FAISS index require significant RAM and CPU resources; not easily scalable to very large datasets without optimization.

## 8. Fairness & Ethics Considerations

### Potential Biases

- **Dataset representation bias:** Caltech-101 contains categories mostly from objects and scenes common in Western contexts; some object types or cultural variations may be underrepresented.
- **Model performance disparity:** SVM classification and retrieval accuracy may vary across underrepresented categories, leading to inconsistent results.

### Ethical Implications

- **Privacy concerns:** If the system were applied to personal or sensitive image data, it could inadvertently expose private information.
- **Potential misuse:** High-accuracy retrieval models could be misused for surveillance or unauthorized content tracking if applied to real-world datasets.

### Mitigation Strategies

- **Dataset curation:** Ensure diverse and representative datasets for broader generalization.

- **Transparency & evaluation:** Regularly audit model performance across different categories to detect bias.
- **Access control:** Restrict usage of the system to ethical, research-oriented contexts; anonymize sensitive data if integrated into practical applications.

## 9. Setup & Run Instructions

### 9.1 Environment Setup

#### Option 1: Using Conda

```
conda create -n project_env python=3.9
conda activate project_env
pip install -r requirements.txt
```

#### Option 2: Using pip / Virtualenv

```
python -m venv venv
# On Linux / Mac:
source venv/bin/activate
# On Windows:
venv\Scripts\activate
pip install -r requirements.txt
```

### 9.2 Hardware Requirements

- **Minimum:** 8GB RAM, CPU only
- **Recommended:** 16GB RAM, NVIDIA GPU (6GB+ VRAM) for faster feature extraction and training

### 9.3 Project Folder Structure Overview

```
image-clustering-retrieval/
├── configs/          # Dataset metadata and config files
├── data/             # Raw and processed images
├── features/         # Extracted embeddings and labels
├── models/           # Saved models and FAISS index
│   ├── classifier/   # SVM, Random Forest, PCA, scaler
│   ├── clustering/   # KMeans, Hierarchical clustering models & metrics
│   └── faiss_index/  # FAISS index for image retrieval
└── results/          # Visualization outputs
└── final_report.json # JSON summary of final evaluation
└── performance_summary.png
```

### 9.4 Running the Project

#### Step 1: Download Dataset

#### Step 2: Preprocess Data

#### Step 3: Extract Features

#### Step 4: Train Models

- **Classification (SVM / Random Forest):**

```
python train_classifier.py --features features/embeddings.npy --labels features/labels.npy
```

- **Clustering (Hierarchical / KMeans):**

```
python train_clustering.py --features features/embeddings.npy --output models/clustering/
```

## Step 5: Build FAISS Index

```
python build_faiss_index.py --features features/embeddings.npy --labels features/labels.npy --output models/faiss_index/
```

## Step 6: Evaluate Model

```
python evaluate.py --model_path models/classifier/svm_rbf.pkl
```

## Step 7: Run Inference / Retrieval (Optional)

```
python predict.py --input sample.jpg --model models/classifier/svm_rbf.pkl  
python retrieve.py --query sample.jpg --index models/faiss_index/image_index.faiss
```

# 10. Conclusion & Future Work

## Summary

In this project, we built a comprehensive **image clustering, retrieval, and classification system** using the Caltech-101 dataset. Key accomplishments include:

- Extracted high-dimensional feature embeddings using **ResNet-50**.
- Applied **hierarchical clustering** to organize images into meaningful groups.
- Implemented **SVM classification** achieving ~86% validation accuracy.
- Developed a **FAISS-based retrieval system** with near-perfect precision@10 for similar image search.
- Generated a complete **evaluation report** including clustering metrics, classification results, and retrieval performance.
- Create a user-friendly interface for real-time image retrieval and classification.

Overall, the system demonstrates strong retrieval capabilities and reliable classification, though clustering quality and dataset diversity present room for improvement.

## Future Improvements

- **Experiment with advanced architectures:** Explore vision transformers (ViT) or EfficientNet for stronger feature extraction.
- **Increase dataset diversity:** Include more categories and images to improve generalization and reduce bias.
- **Improve clustering methods:** Try deep clustering or spectral clustering to better separate overlapping categories.
- **Optimize computational efficiency:** Implement dimensionality reduction or approximate nearest neighbor search for large-scale datasets.
- **Bias mitigation and fairness:** Continuously monitor performance across categories and ensure ethical use of the system.